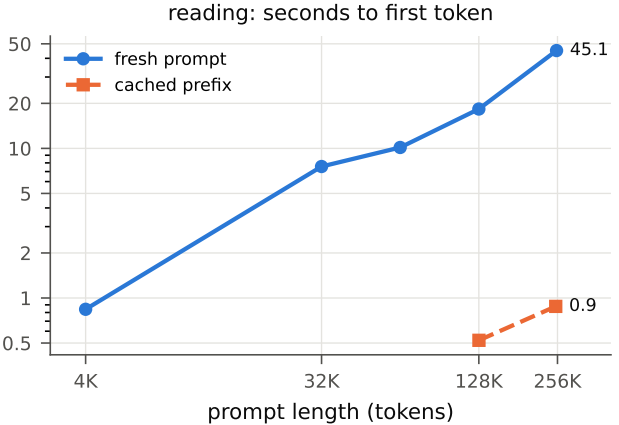
reading: seconds to first token
fresh prompt
cached prefix
50
20
10
5
2
1
0.5
45.1
0.9
4K
32K
128K
256K
prompt length (tokens)

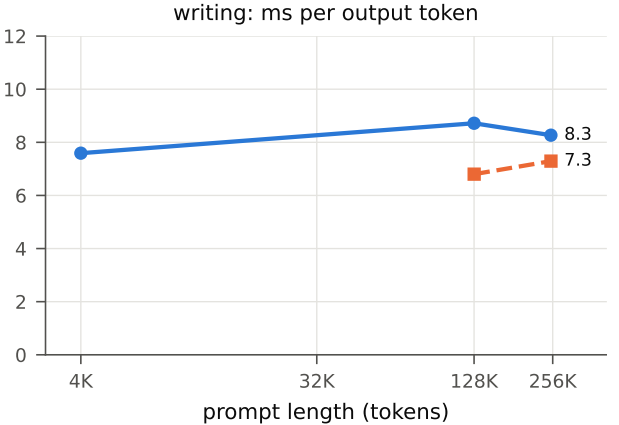
writing: ms per output token
12
10
8
6
4
2
0
8.3
7.3
4K
32K
128K
256K
prompt length (tokens)