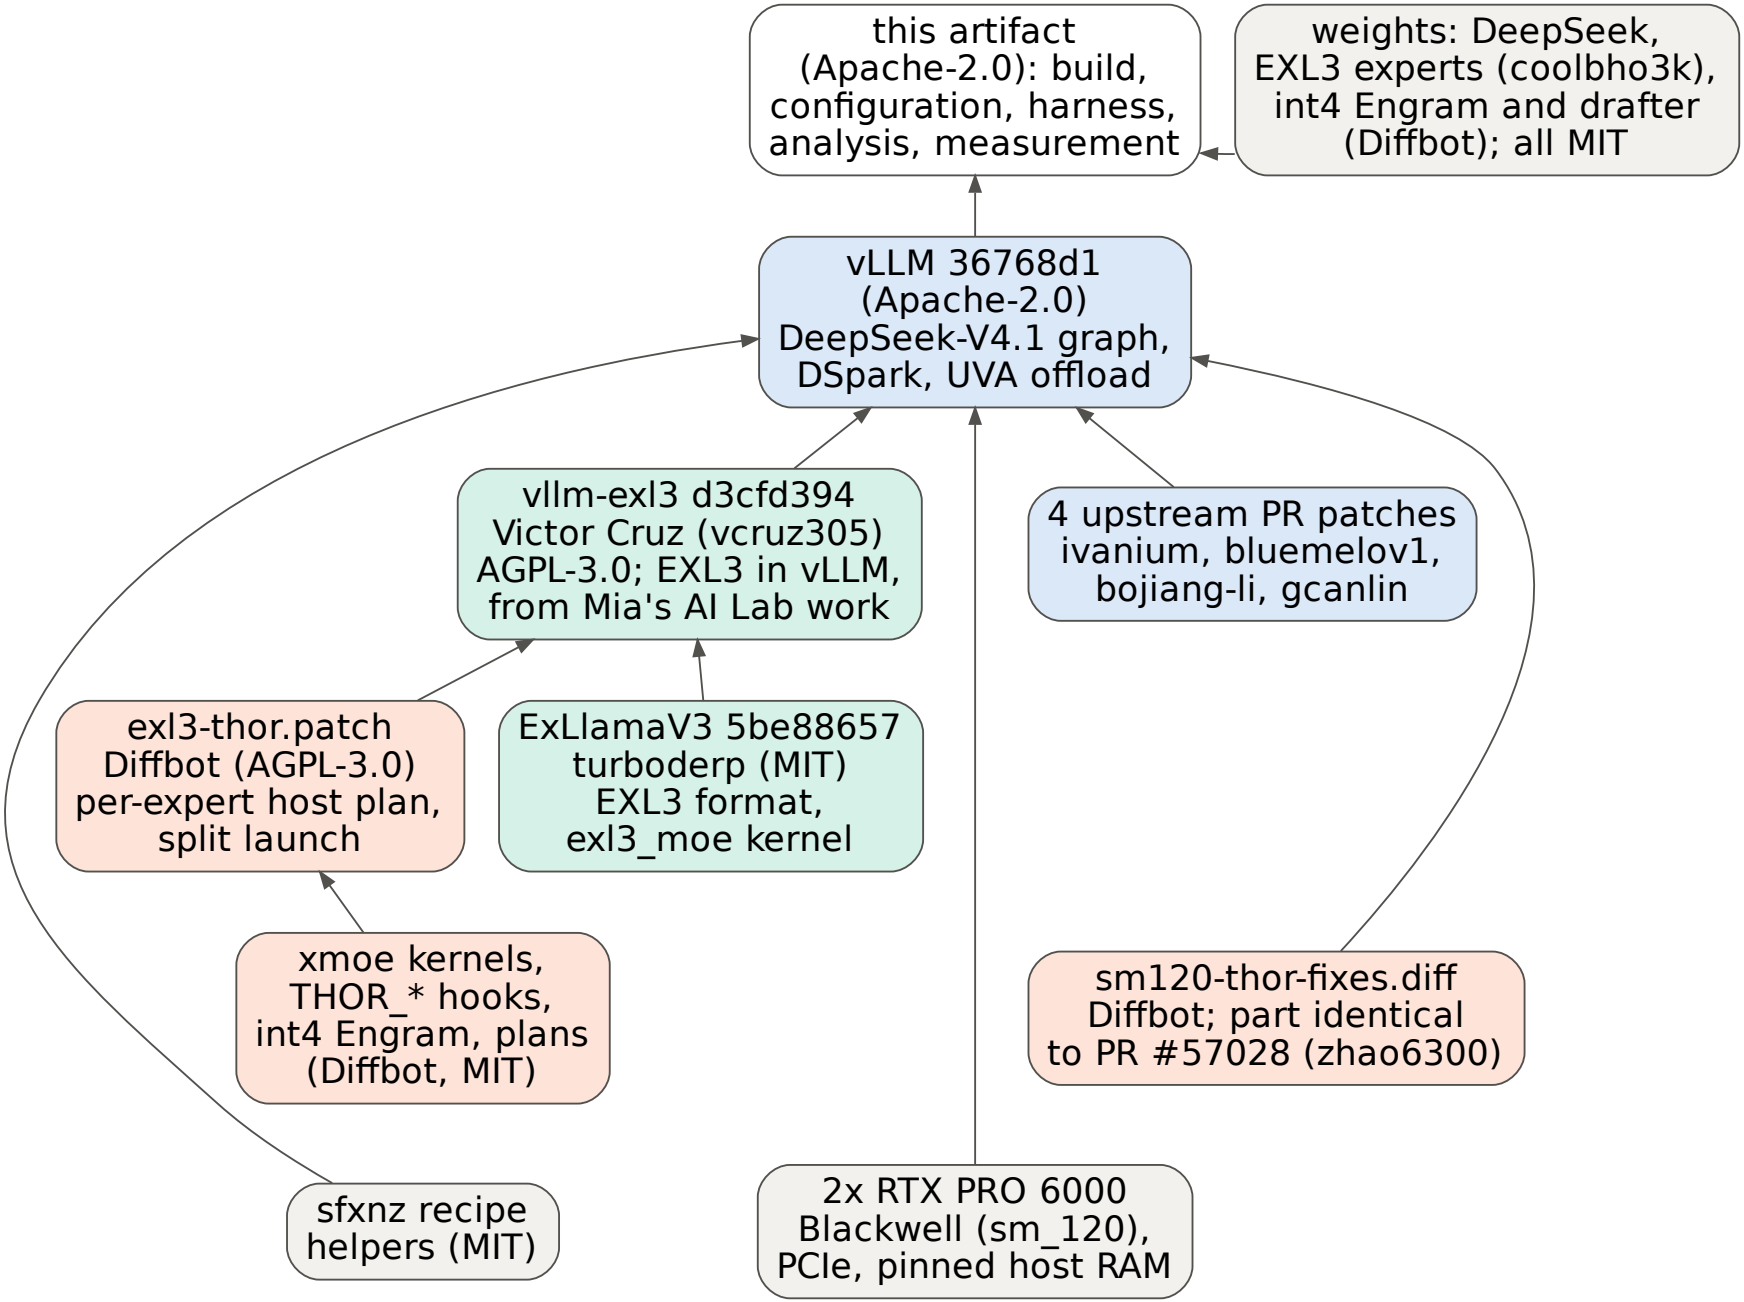
this artifact
(Apache-2.0): build,
configuration, harness,
analysis, measurement
weights: DeepSeek,
EXL3 experts (coolbho3k),
int4 Engram and drafter
(Diffbot); all MIT
vLLM 36768d1
(Apache-2.0)
DeepSeek-V4.1 graph,
DSpark, UVA offload
vllm-exl3 d3cfd394
Victor Cruz (vcruz305)
AGPL-3.0; EXL3 in vLLM,
from Mia's AI Lab work
4 upstream PR patches
ivanium, bluemelov1,
bojiang-li, gcanlin
exl3-thor.patch
Diffbot (AGPL-3.0)
per-expert host plan,
split launch
ExLlamaV3 5be88657
turboderp (MIT)
EXL3 format,
exl3_moe kernel
xmoe kernels,
THOR_* hooks,
int4 Engram, plans
(Diffbot, MIT)
sm120-thor-fixes.diff
Diffbot; part identical
to PR #57028 (zhao6300)
sfxnz recipe
helpers (MIT)
2x RTX PRO 6000
Blackwell (sm_120),
PCIe, pinned host RAM